



MEDPREDICT

SHAHD SALAH

DEPI_FINAL PROJECT

TABLE OF CONTENTS

INTRODUCTION

DATA LINKS

THE PROCESS

1-BREAST CANCER


2-DIABETES

3-STROKE

WHAT IS IN THE CODE

CONCLUSION

INTRODUCTION




I HAVE DEVELOPED THREE SEPARATE MACHINE LEARNING-POWERED WEBSITES DESIGNED TO PREDICT **STROKE**, **DIABETES**, AND **BREAST CANCER**. EACH WEBSITE INTEGRATES USER-FRIENDLY INTERFACES BUILT WITH **GRADIO**, ALLOWING USERS TO INPUT THEIR PERSONAL HEALTH DATA AND RECEIVE A PREDICTION FOR EACH DISEASE. THE MODELS BEHIND THE SITES EMPLOY ADVANCED TECHNIQUES LIKE ENSEMBLE METHODS AND STACKING CLASSIFIERS TO ENSURE ACCURACY AND RELIABILITY.

these websites allow users to input relevant medical and lifestyle data, offering predictions based on machine learning models that include ensemble methods and deep learning architectures. Each model has been carefully trained and fine-tuned to provide high accuracy, making the applications practical tools for preliminary health risk assessments. The user interface is designed to be intuitive, ensuring accessibility to a wide audience while maintaining robust predictive performance.



DATA LINKS



diabetes.csv

Kaggle is the world's largest data science community with powerful tools and resources to help you achieve your data science goals.

[k kaggle.com](https://www.kaggle.com)



Breast Cancer Dataset

Explore and run machine learning code with Kaggle Notebooks | Using data from Breast Cancer Dataset

[k Kaggle](https://www.kaggle.com) / Mar 20



Stroke Prediction Dataset

11 clinical features for predicting stroke events

[k kaggle.com](https://www.kaggle.com)



THE PROCESS



**DATA
PREPROCESSING**



MODELS USED



INTERFACE

I-BREAST CANCER

ADD LIBRARIES AND FILE

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, precision_recall_fscore_support, recall_score, precision_score, f1_score, ConfusionMatrixDisplay
from sklearn.linear_model import LogisticRegression
import gradio as gr
import warnings
warnings.filterwarnings('ignore')
```

```
data=pd.read_csv('Breast_cancer.csv')
data.head()
```

id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	...	texture_worst	perimeter_worst
842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	...	17.33	184.60
842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	...	23.41	158.60
84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	...	25.53	152.50
84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28360	0.2414	0.10530	...	26.50	98.50

2-ANALYZE DATA USING INFO, DESCRIPTION, VALUE_COUNTS AND DUPLICATED

```
data.describe()
```

✓ 0.1s

Python

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	symmetry_mean	fractal_dimension_mean	...	radius_1
count	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	...	569.00
mean	14.127292	19.289649	91.969033	654.889104	0.096360	0.104341	0.088799	0.048919	0.181162	0.062798	...	16.26
std	3.524049	4.301036	24.298981	351.914129	0.014064	0.052813	0.079720	0.038803	0.027414	0.007060	...	4.82
min	6.981000	9.710000	43.790000	143.500000	0.052630	0.019380	0.000000	0.000000	0.106000	0.049960	...	7.92
25%	11.700000	16.170000	75.170000	420.300000	0.086370	0.064920	0.029560	0.020310	0.161900	0.057700	...	13.01
50%	13.370000	18.840000	86.240000	551.100000	0.095870	0.092630	0.061540	0.033500	0.179200	0.061540	...	14.97
75%	15.780000	21.800000	104.100000	782.700000	0.105300	0.130400	0.130700	0.074000	0.195700	0.066120	...	18.75
max	28.110000	39.280000	188.500000	2501.000000	0.163400	0.345400	0.426800	0.201200	0.304000	0.097440	...	36.04

8 rows × 30 columns

```
data.duplicated().sum()
```

✓ 0.0s

Python

0

+ Code

+ Markdown

```
data['diagnosis'].value_counts()
```

✓ 0.0s

Python

diagnosis

B 357

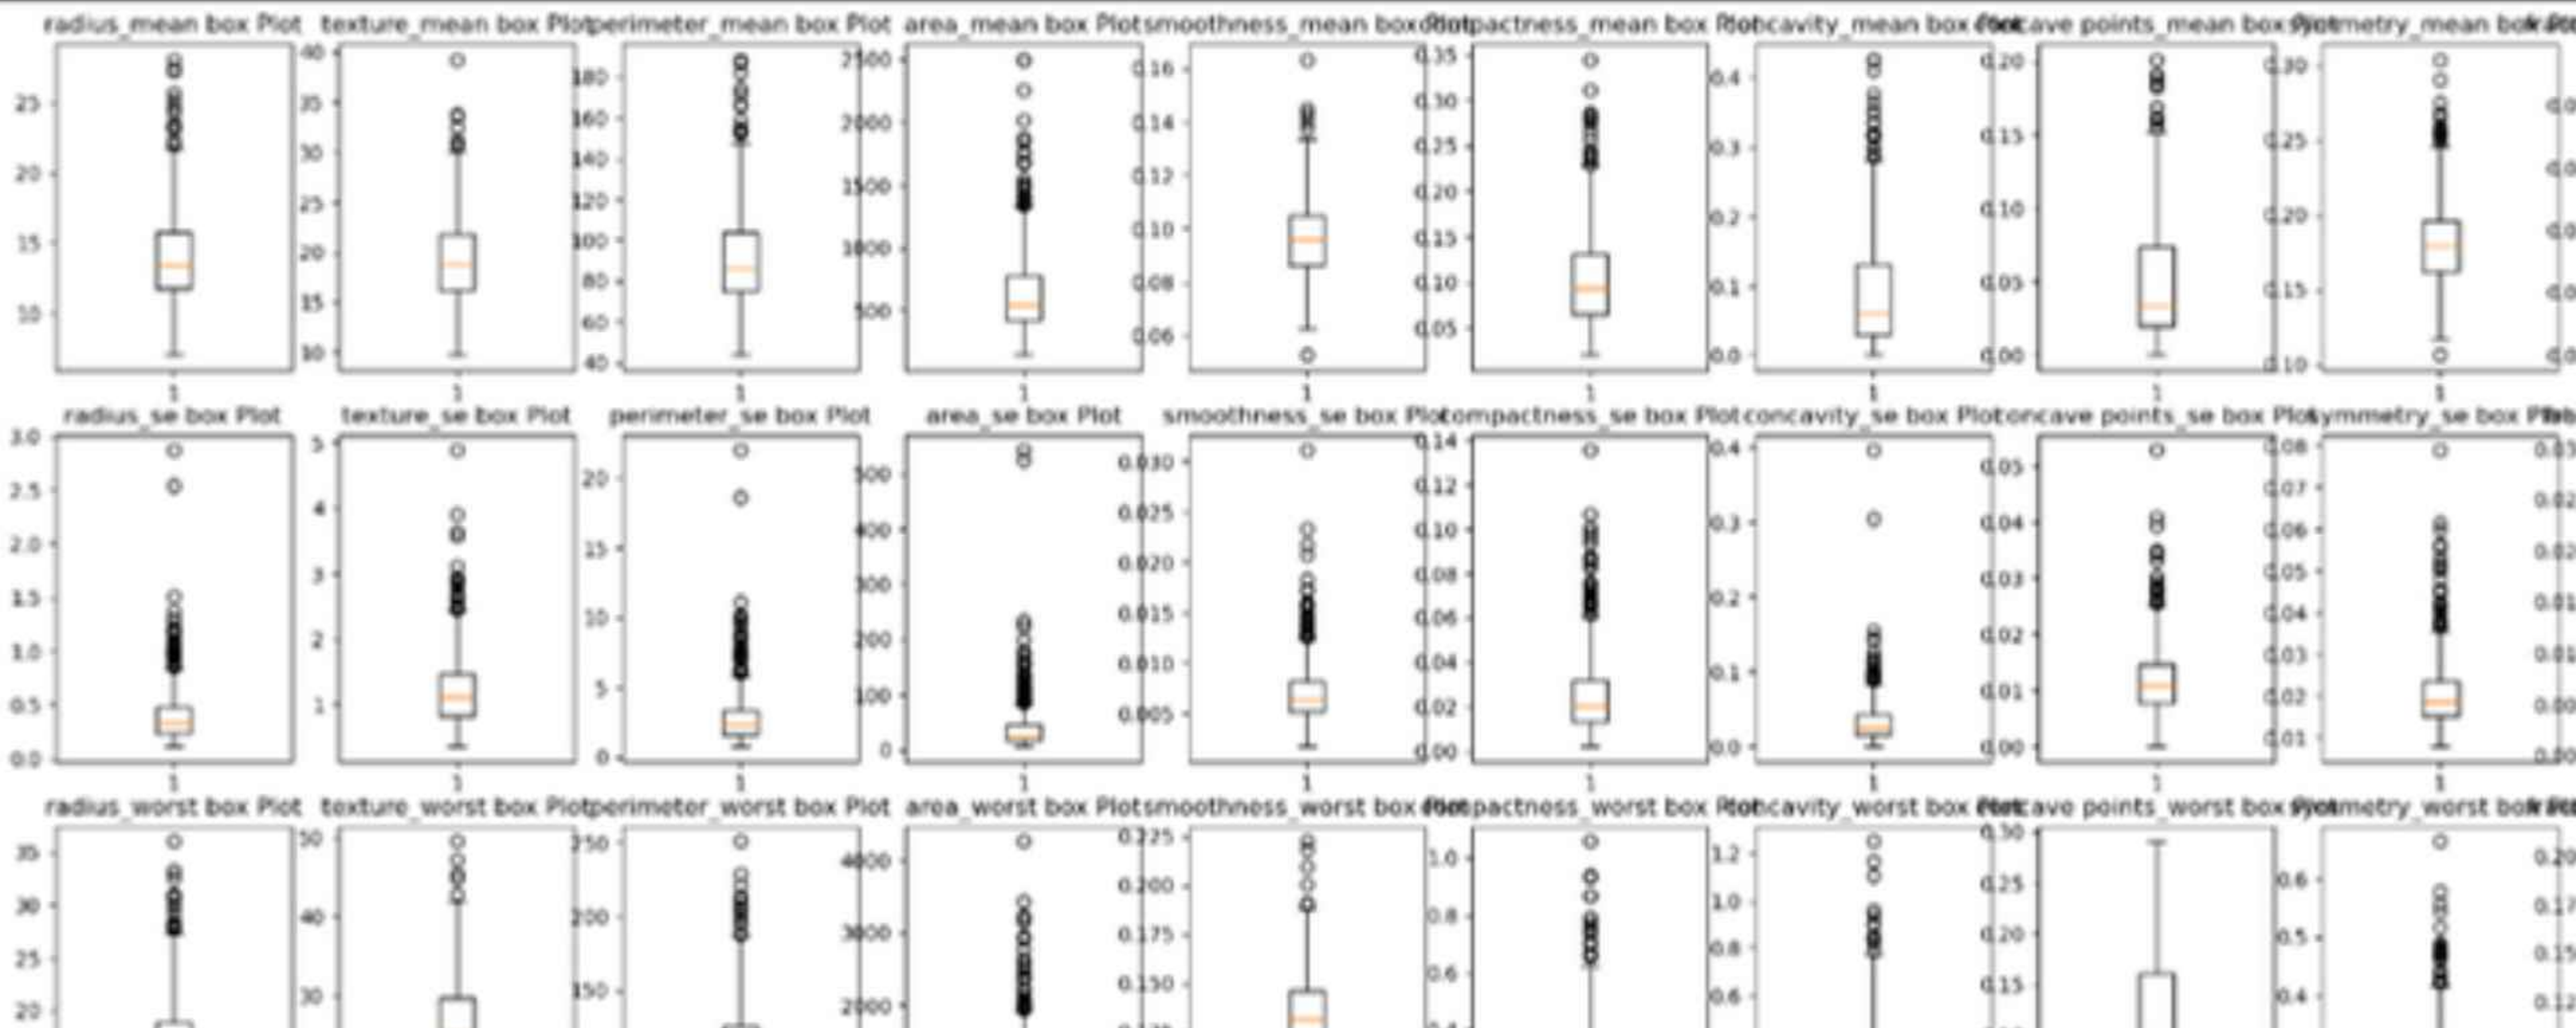
M 212

3-LABEL ENCODE TO BE 0,1

```
labelencoder = LabelEncoder()
data['diagnosis'] = labelencoder.fit_transform(data['diagnosis'])
```

4-BOX PLOT

```
numCols = data.select_dtypes('float64').columns
numCols
plt.figure(figsize=(25,10))
for i, col in enumerate(numCols):
    plt.subplot(3, 10, i+1)
    plt.boxplot(data[col])
    plt.title(f"{col} box Plot")
```



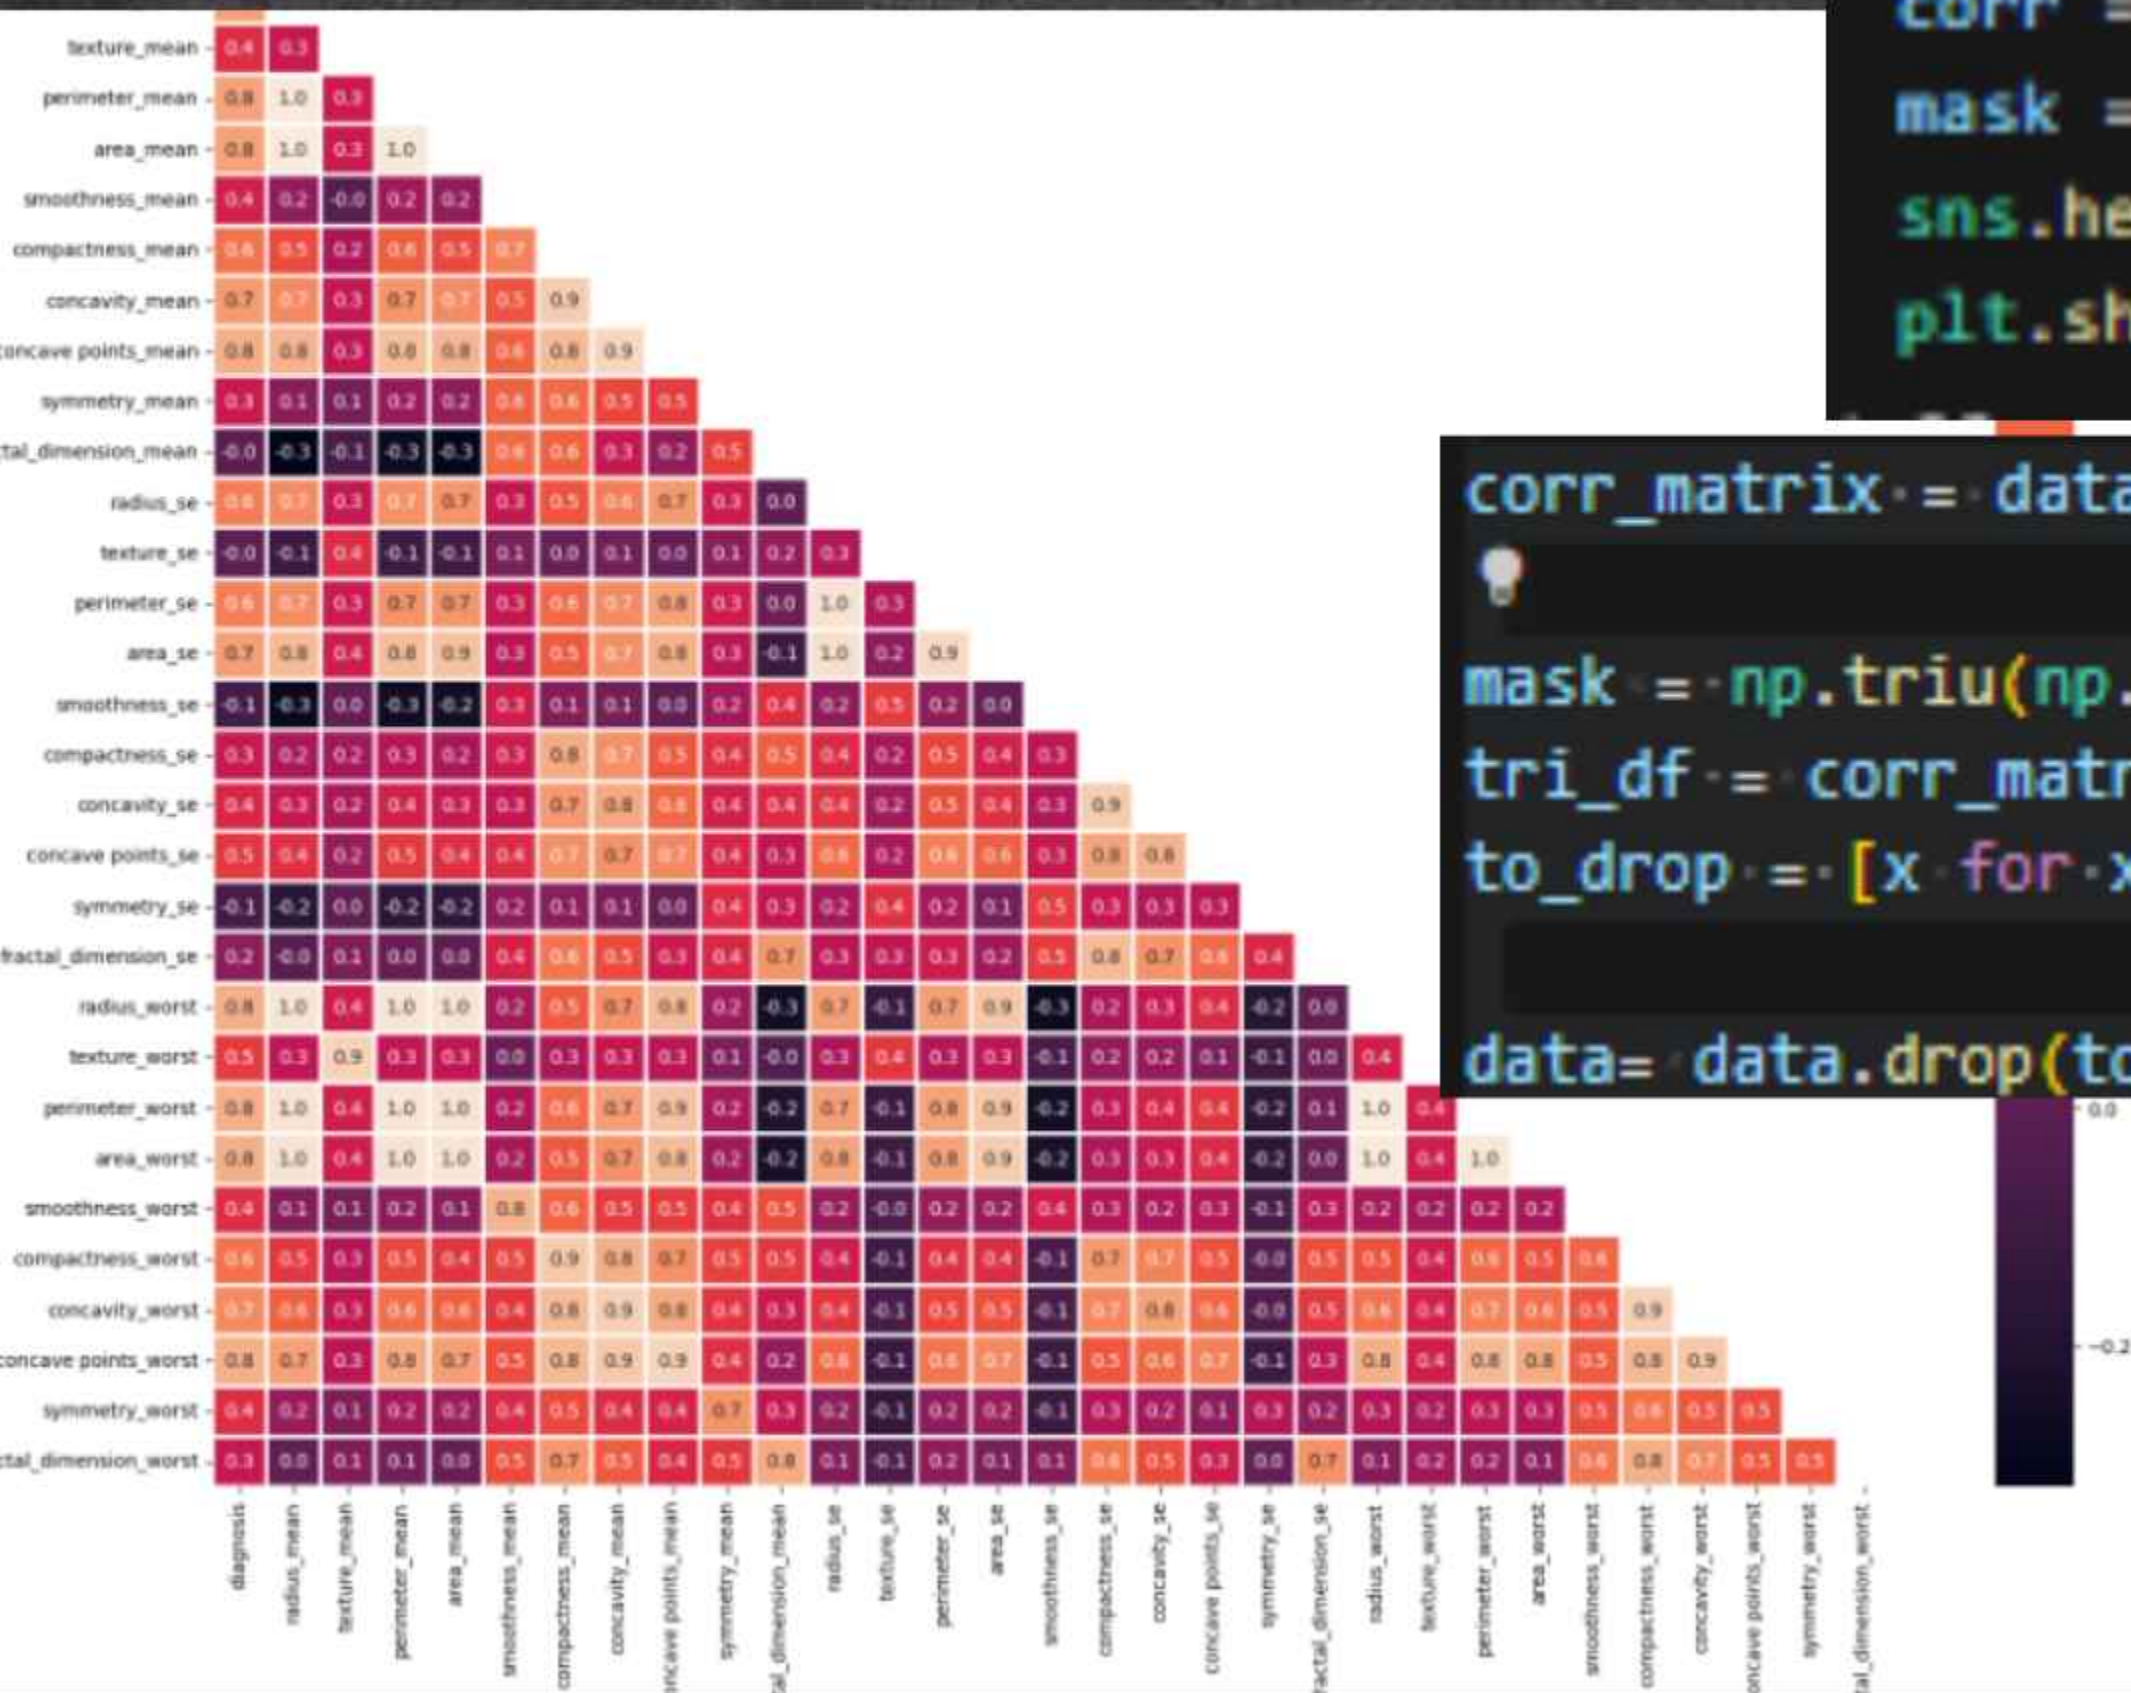
5-SOLVEING THE PROBLEM OF OUT LAYERS

```
for col in numCols:
    Q1 = data[col].quantile(0.25)
    Q3 = data[col].quantile(0.75)
    IQR = Q3 - Q1
    Lf = Q1 - 1.5 * IQR
    Up = Q3 + 1.5 * IQR

    # Cap values below the lower fence and above the upper fence
    data[col] = data[col].clip(lower=Lf, upper=Up)
```

✓ 0.1s

6-DELETING SOME FEATURES



```
plt.figure(figsize=(20,15))
corr = data.corr()
mask = np.triu(np.ones_like(corr, dtype=bool))
sns.heatmap(corr, mask=mask, linewidths=1, annot=True,
            plt.show()
```

```
corr_matrix = data.corr().abs()
mask = np.triu(np.ones_like(corr_matrix, dtype=bool))
tri_df = corr_matrix.mask(mask)
to_drop = [x for x in tri_df.columns if any(tri_df[x] > 0.90)]
data = data.drop(to_drop, axis=1)
```


7- FINAL TOUCHES BEFORE MODELS

```
x = data.drop(columns=['diagnosis'])  
y = data.diagnosis  
x.shape, y.shape
```

✓ 0.0s

((569, 19), (569,))

```
from imblearn.over_sampling import SMOTE  
smote = SMOTE()  
X_resampled, y_resampled = smote.fit_resample(X, y)  
X_resampled.shape, y_resampled.shape
```

✓ 0.8s

((714, 19), (714,))

```
X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled, test_size=0.3, random_state=42)
```

✓ 0.0s

+ Code

+ Markdown

```
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

✓ 0.0s

((499, 19), (215, 19), (499,), (215,))

```
scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)
```


✓ 0.0s


NOTE :
SMOTE (SYNTHETIC
MINORITY OVER-
SAMPLING TECHNIQUE)

MODELS ANALYSIS

```
model_names = ['LogisticRegression', 'Random Forest', 'KNeighbors']
accuracies = [acc , acc_RF, auc_score, ]

plt.figure(figsize=(6, 9))
plt.bar(model_names, accuracies, color='skyblue')
plt.title('Comparison of Model Test Accuracies')
plt.xlabel('Model')
plt.ylabel('Test Accuracy')
for i, accuracy in enumerate(accuracies):
    plt.text(i, accuracy + 0.005, f'{accuracy:.4f}', ha='center', va='bottom')
plt.ylim(0.8, 1)
plt.show()
```





output

Malignant

Flag



output

Benign

Flag

2-DIABETES


PROCESS ZERO VALUES

```
#for zeros
def rep(data,z):
    med=data[z].median()
    data[z]=data[z].replace(0,med)
    return data
```

```
data=rep(data,'BMI')
data=rep(data,'SkinThickness')
data=rep(data,'Insulin')
data=rep(data,'Pregnancies')
data=rep(data,'Glucose')
data=rep(data,'BloodPressure')
```

VISUALIZE

```
numCols = data.select_dtypes('number').columns.tolist()
numCols.remove('Outcome')
for col in numCols:
    plt.hist(data[data['Outcome']==0][col],10,label='NON DIABETES')
    plt.hist(data[data['Outcome']==1][col],10,label='DIABETES')
    plt.legend(col)
    plt.ylabel('Freq.')
    plt.title(f'Histogram of {col}')
    plt.show()
```



LINEAR REGRESSION

LINEAR REGRESSION MODEL

```
X=data.drop('DiabetesPedigreeFunction', axis=1)
y=data['DiabetesPedigreeFunction']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
models = {
    'Linear Regression': LinearRegression(),
    'Ridge Regression': Ridge(),
    'Lasso Regression': Lasso(),
    'Random Forest': RandomForestRegressor(n_estimators=100, random_state=42)
}
```


(Fit models and evaluate)

```
for model_name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    mse = mean_squared_error(y_test, y_pred)
    print(f"{model_name} MSE: {mse:.4f}")
```

Linear Regression MSE: 0.0856
Ridge Regression MSE: 0.0856
Lasso Regression MSE: 0.0951
Random Forest MSE: 0.1017

```
# Cross-validation for better estimation
for model_name, model in models.items():
    cv_scores = cross_val_score(model, X_train, y_train, scoring='neg_mean_squared_error')
    print(f"{model_name} Cross-validated MSE: {-np.mean(cv_scores):.4f}")
```

Linear Regression Cross-validated MSE: 0.1108
Ridge Regression Cross-validated MSE: 0.1108
Lasso Regression Cross-validated MSE: 0.1138
Random Forest Cross-validated MSE: 0.1172

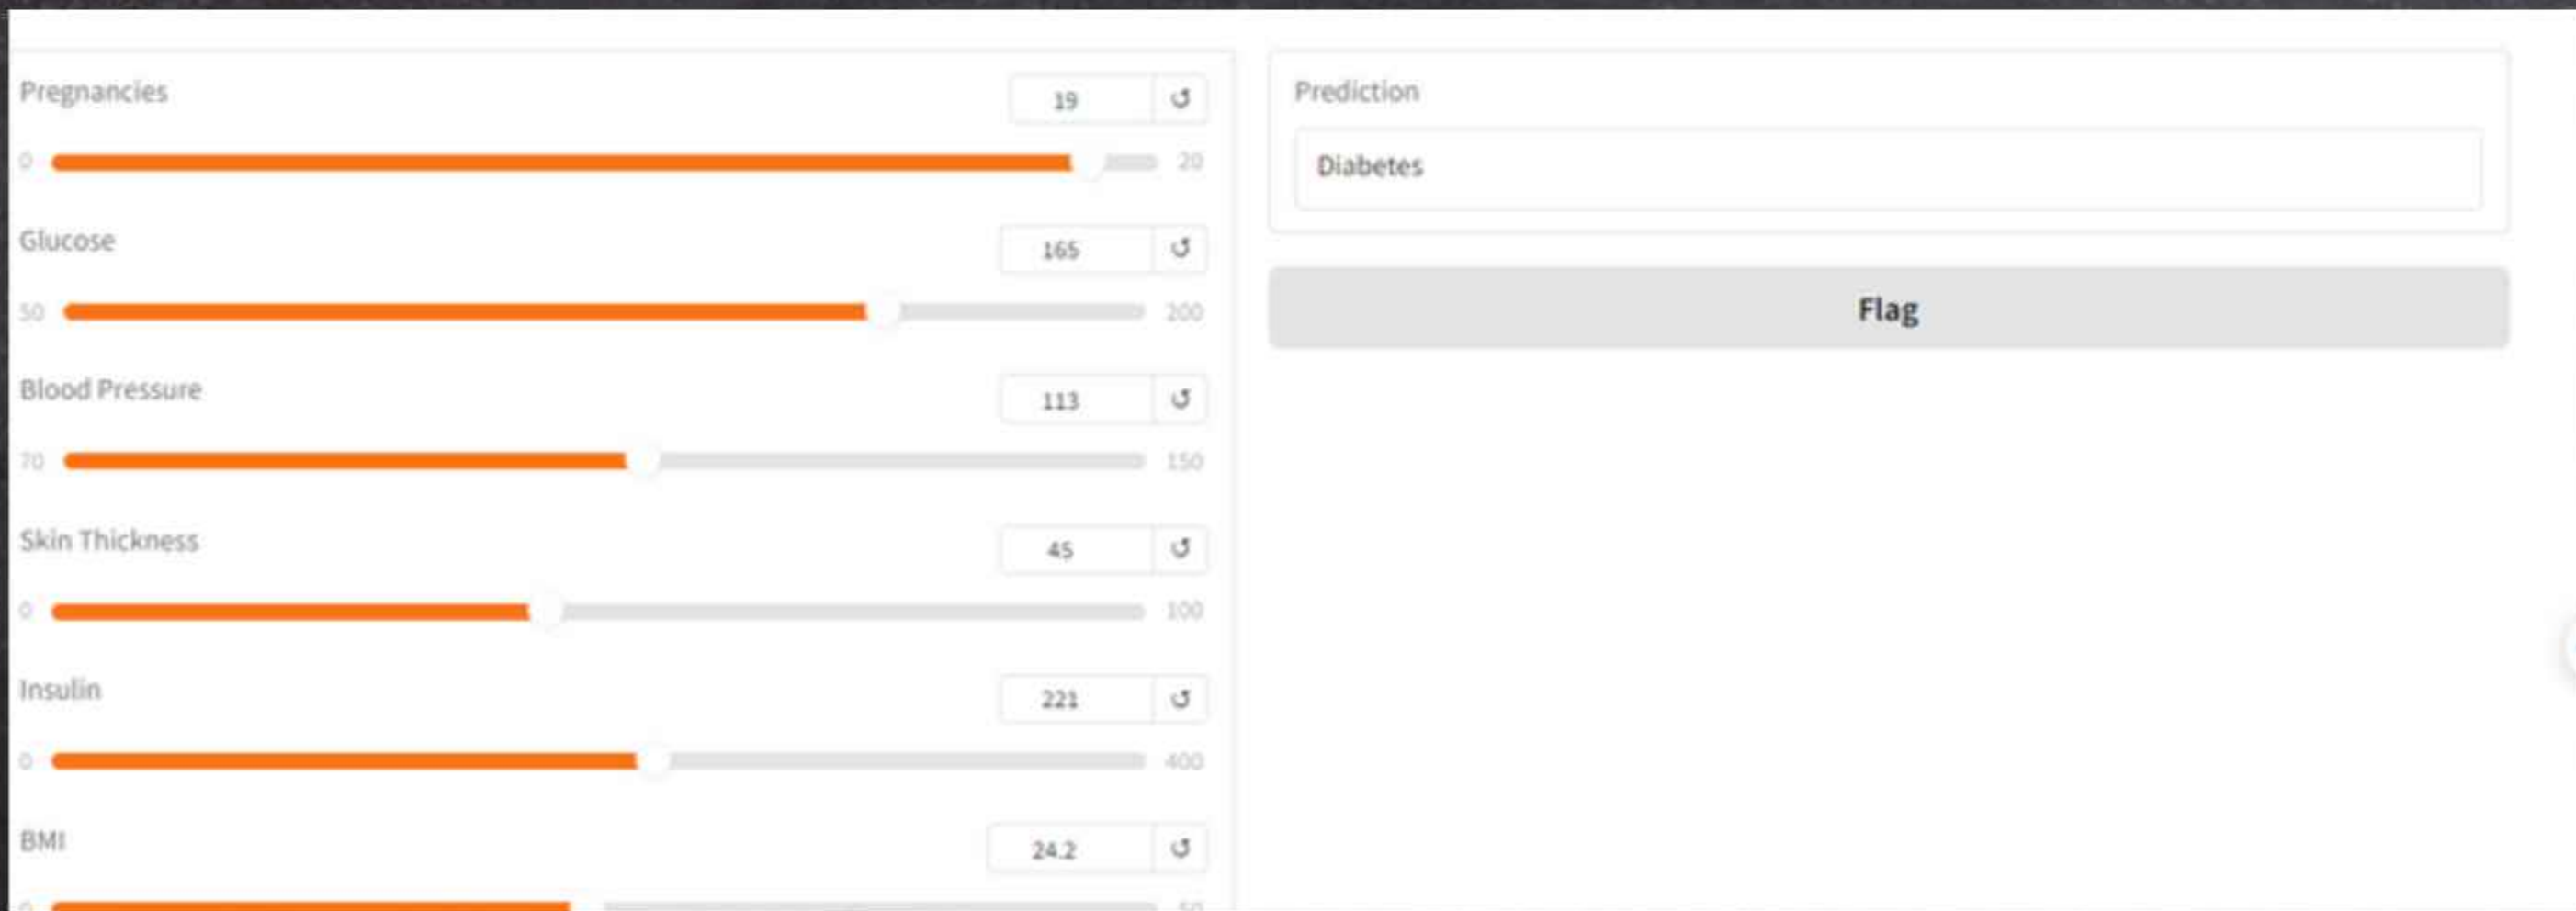
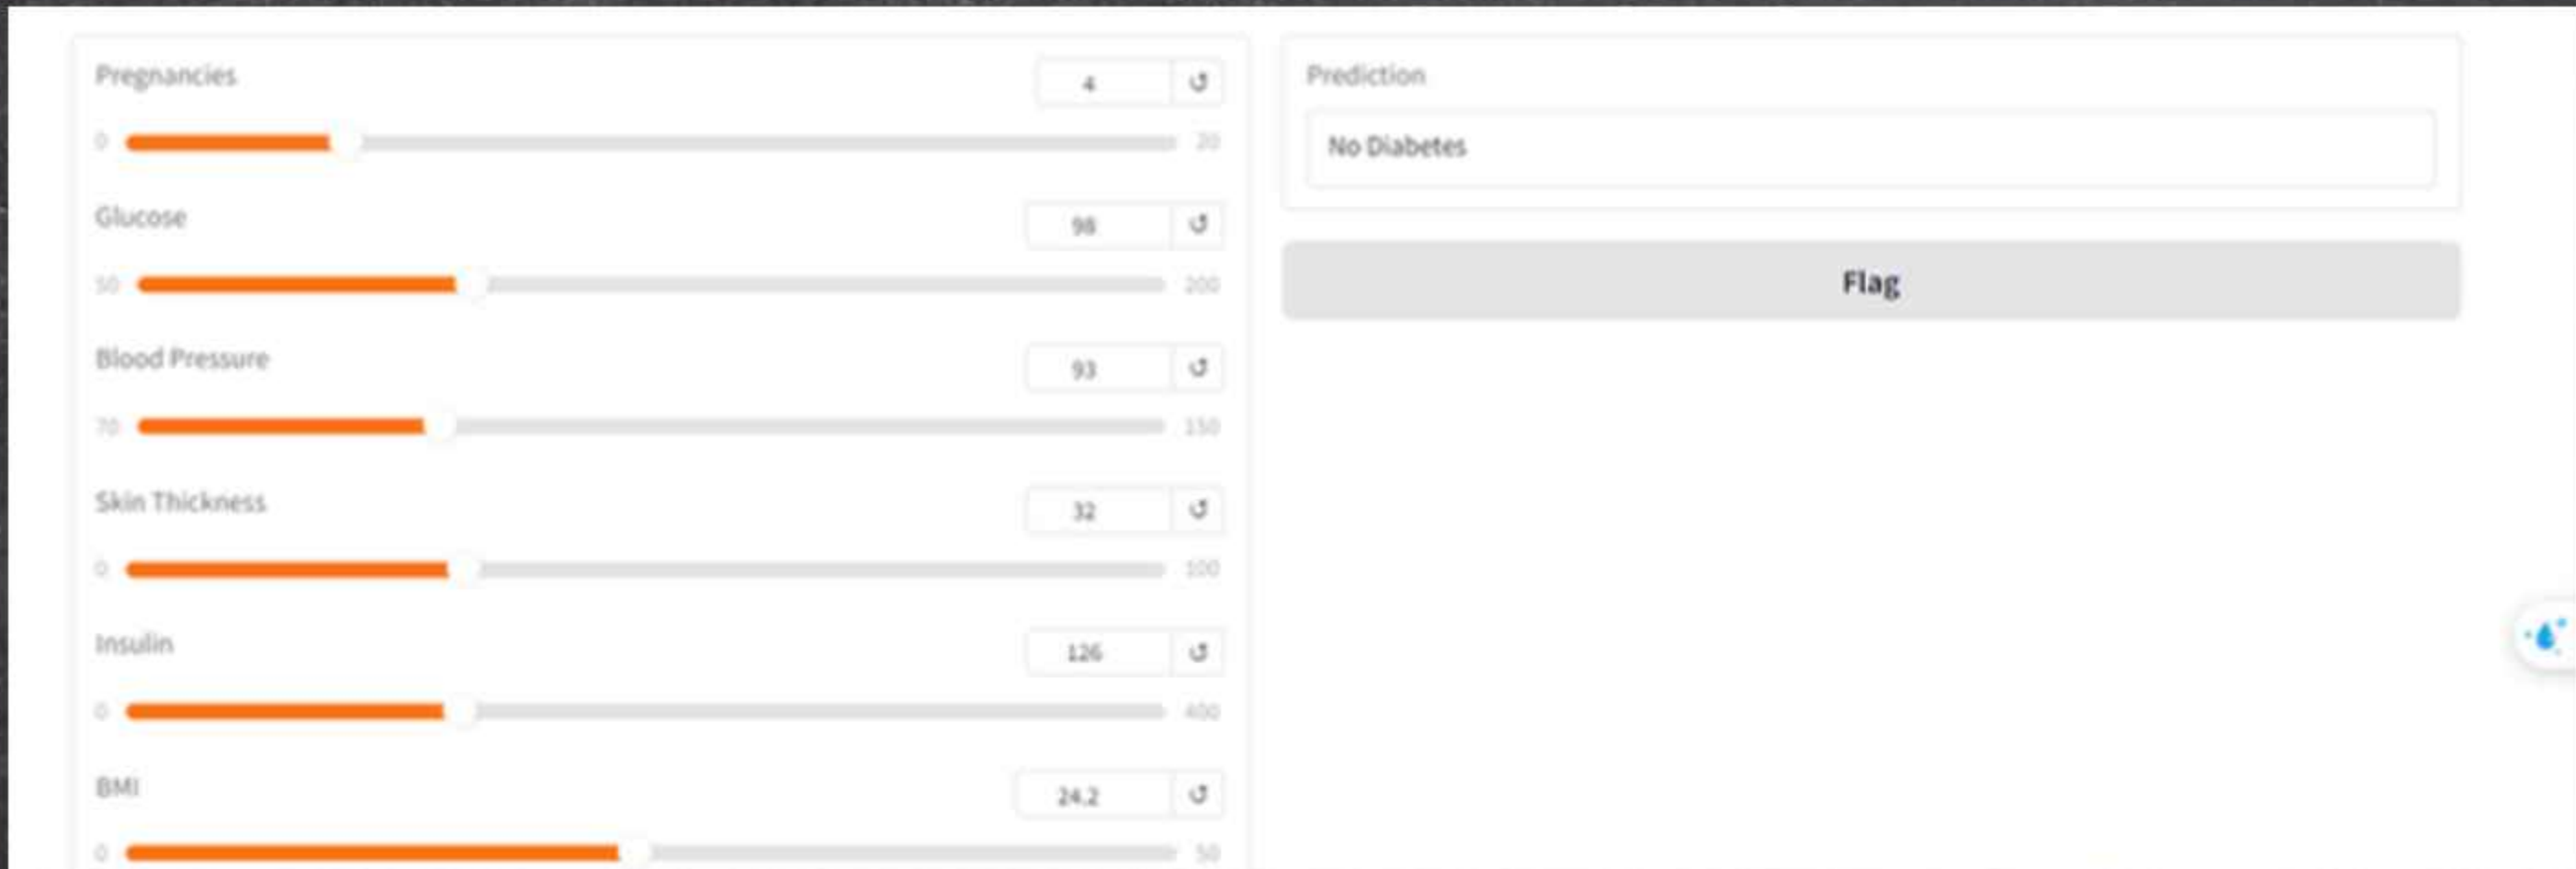
```
from sklearn.preprocessing import PolynomialFeatures

poly = PolynomialFeatures(degree=3, interaction_only=False)
X_poly = poly.fit_transform(X_train)
X_test_poly = poly.transform(X_test)

model_poly = LinearRegression()
model_poly.fit(X_poly, y_train)
y_pred_poly = model_poly.predict(X_test_poly)
mse_poly = mean_squared_error(y_test, y_pred_poly)
print(f"Polynomial Regression MSE: {mse_poly}")
```

LAYERS

```
model.add(Dense(256, activation='relu', input_shape=(X_train_scaled.shape[1],)))
model.add(Dropout(0.5)) # Adding dropout for regularization
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid')) # Output layer for binary classification
```

3-STROKE

```
data['stroke'].value_counts()  
###SMOTE TARGET
```

```
stroke  
0    4861  
1     249  
Name: count, dtype: int64
```

```
print("skew is =" ,data['bmi'].skew())  
pd.DataFrame(data[['bmi' , 'avg_glucose_level']]).describe().T
```

```
skew is = 1.0553402052962912
```

	count	mean	std	min	25%	50%	75%	max
bmi	4909.0	28.893237	7.854067	10.30	23.500	28.100	33.10	97.60
avg_glucose_level	5110.0	106.147677	45.283560	55.12	77.245	91.885	114.09	271.74

```
data_na=data.loc[data['bmi'].isnull()]  
data_na['stroke'].value_counts()
```

```
stroke  
0    161  
1     40  
Name: count, dtype: int64
```

so we cant drop


```
data['bmi'].fillna(data['bmi'].median(), inplace=True)  
pd.DataFrame(data['bmi']).describe().T
```



```
# One-hot encode the categorical features
data = pd.get_dummies(data, columns=['gender', 'smoking_status', 'work_type'])
data.info()
```

```
bmi_bins = [0, 18.5, 24.9, 29.9, 34.9, 100]
bmi_labels = ['Underweight', 'Normal', 'Overweight', 'Obese', 'Severely Obese']
data['bmi_category'] = pd.cut(data['bmi'], bins=bmi_bins, labels=bmi_labels)
bins = [0, 18, 35, 50, 65, 100]
labels = ['Child', 'Young Adult', 'Adult', 'Senior', 'Elderly']
data['age_group'] = pd.cut(data['age'], bins=bins, labels=labels)
data
```

```
# Map age_group and bmi_category to consistent numerical values
age_mapping = {'Child': 1, 'Young Adult': 2, 'Adult': 3, 'Senior': 4, 'Elderly': 5}
bmi_mapping = {'Underweight': 1, 'Normal': 2, 'Overweight': 3, 'Obese': 4, 'Severely Obese': 5}
data['age_group'].replace(age_mapping, inplace=True)
data['bmi_category'].replace(bmi_mapping, inplace=True)
data
```


```
from imblearn.over_sampling import RandomOverSampler
ov=RandomOverSampler(sampling_strategy='minority')
X = data.drop(columns=['stroke'])
y = data['stroke']
X_ov , y_ov = ov.fit_resample(X,y)
```

```
# Create a stacking model
stacking_model = StackingClassifier(
    estimators=[('logreg', model), ('rf', rf_clf)],
    final_estimator=LogisticRegression()
)

# Fit the stacking model
stacking_model.fit(X_train, y_train)

# Predictions
y_pred_stacking = stacking_model.predict(X_test)
```



```

def calculate_bmi_category(bmi):
    if bmi < 18.5:
        return 1 # Underweight
    elif 18.5 <= bmi < 24.9:
        return 2 # Normal
    elif 25 <= bmi < 29.9:
        return 3 # Overweight
    elif 30 <= bmi < 34.9:
        return 4 # Obese
    else:
        return 5 # Severely Obese

# Function to calculate age group
def calculate_age_group(age):
    if age < 18:
        return 1 # Child
    elif 18 <= age < 35:
        return 2 # Young Adult
    elif 35 <= age < 55:
        return 3 # Adult
    elif 55 <= age < 75:
        return 4 # Senior
    else:
        return 5 # Elderly

# Define a function to preprocess the inputs and predict stroke
def predict_stroke(gender, age, hypertension, heart_disease, marital_status, work_type,
                  residence_type, avg_glucose_level, bmi, smoking_status):

    # Convert categorical inputs to numeric
    hypertension = 1 if hypertension else 0
    heart_disease = 1 if heart_disease else 0
    marital_status = 1 if marital_status == "Yes" else 0

    inputs = [
        gr.Radio(choices=['Male', 'Female'], label="Gender"),
        gr.Slider(0, 100, label="Age"),
        gr.Checkbox(label="Hypertension"),
        gr.Checkbox(label="Heart Disease"),
        gr.Radio(choices=['Yes', 'No'], label="Ever Married"),
        gr.Dropdown(label="Work Type", choices=["Private", "Self-employed", "Govt-job", "Children", "Never worked"]),
        gr.Radio(label="Residence Type", choices=["Urban", "Rural"]),
        gr.Slider(30, 300, label="Average Glucose Level"),
        gr.Slider(5, 45, label="BMI"),
        gr.Dropdown(label="Smoking Status", choices=["Never smoked", "Formerly smoked", "Smokes", "Unknown"])
    ]

    output = gr.Textbox(label="Prediction")

    # Create Gradio interface
    gr.Interface(fn=predict_stroke, inputs=inputs, outputs=output, title="Stroke Prediction",
                description="Enter patient details to predict the risk of stroke.").launch()

    from collections import Counter
    print("Original class distribution:", Counter(y_ov))

```


Gender
☐ Male ☒ Female

Age
20

☒ Hypertension
☐ Heart Disease

Ever Married
☒ Yes ☐ No

Work Type
Govt job

Residence Type
☐ Urban ☒ Rural

Average Glucose Level
104

BMI
23.5

Smoking Status
Formerly smoked

Prediction
No Stroke

Flag

Gender
☐ Male ☒ Female

Age
58

☐ Hypertension
☒ Heart Disease

Ever Married
☒ Yes ☐ No

Work Type
Private

Residence Type
☐ Urban ☒ Rural

Average Glucose Level
207


BMI
34

Smoking Status
Smokes

Prediction
Stroke

Flag

WHAT IS IN THE CODE



CONCLUSION

SUMMARIZE THE PRESENTATION

these projects underscore the transformative impact of machine learning in the healthcare domain. By effectively utilizing data-driven approaches, we can develop predictive models that not only improve diagnostic accuracy but also facilitate timely interventions, leading to enhanced patient care. These models serve as valuable tools for healthcare professionals, allowing them to make informed decisions and prioritize patient health more effectively. As we continue to refine these models and explore new data sources, the potential for advancements in disease prediction and prevention remains significant.